

Operating Systems

4. CPU: Process

Hyunchan, Park

<http://oslab.chonbuk.ac.kr>

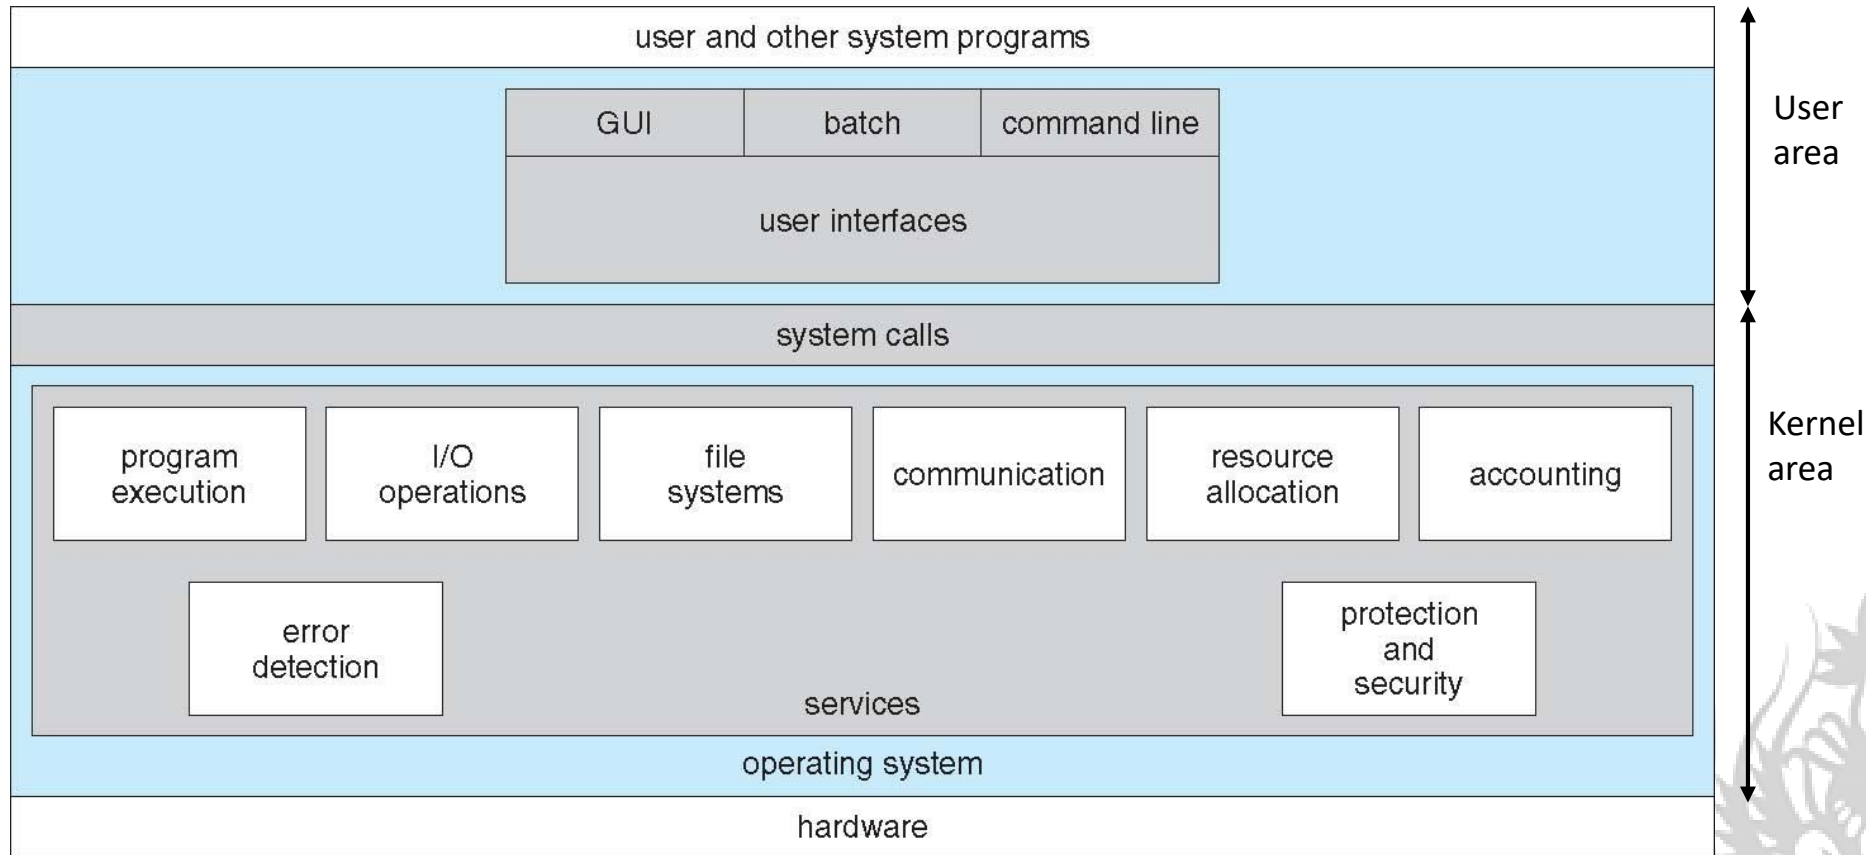
Division of Computer Science and Engineering

Chonbuk National University

Contents

- Process: concept
- Process: state diagram
- PCB: In kernel data-structure of process
- Context switching
- Creation and termination of process

A View of Operating System Services




Scope of Operating System

- 프로세스 관리 및 스케줄링 (CPU)
- 가상 메모리 관리 (Memory)
- 파일 관리 (Secondary Storage)
- I/O 시스템 관리
- 네트워킹
- 보안



프로그램이 만들어지는 과정



프로그램이 만들어지는 과정(cont.)

- 소스코드 (.c file)
 - 프로그램이 수행하고자 하는 작업이 프로그래밍 언어로 표현
- 컴파일러 (Compiler)
 - 사람이 이해할 수 있는 프로그래밍 언어로 작성된 소스 코드를 컴퓨터(CPU)가 이해할 수 있는 기계어로 표현된 Object file로 변환
- 오브젝트 (.o file)
 - 컴퓨터가 이해할 수 있는 기계어로 구성된 파일. 자체로는 수행이 이루어지지 못함. 프로세스로 변환되기 위한 정보가 삽입되어야 함
 - Relocatable addresses (relative addresses) 로 표현
 - 심볼들의 주소가 상대적인 값으로 표현됨
 - 예) 시작 주소로부터 26바이트 지점

프로그램이 만들어지는 과정(cont.)

- 링커 (Linker, Linkage editor)
 - 관련된 여러 오브젝트 파일들과 라이브러리들을 연결하여 메모리로 로드 될 수 있는 하나의 실행파일로 작성
- 실행파일 (.exe와 같은 executable file)
 - 특정한 환경(OS)에서 수행될 수 있는 파일. 프로세스로의 변환을 위한 header, 작업 내용인 text, 필요한 데이터인 data를 포함한다
 - Absolute addresses로 표현
 - 심볼들의 주소가 절대값으로 표현됨
 - 예) 0x0100 3F0A
- 컴파일러와 링커는 결과물이 수행될 OS와 CPU에 따라 다른 형태의 파일을 만든다

Process Concept

- 정의

- A program in execution
- Abstraction for CPU sharing
 - CPU를 공유하는 여러 작업 간의 구분을 위한 단위로 사용
 - 각 프로세스는 컴퓨터 시스템을 독점해서 사용하는 것으로 인식
 - 상호 간의 간섭, 침범 불가

- 구성


- CPU 상태: registers including program counter
- 메모리 (virtual address space)
 - Text section (프로그램 코드 저장)
 - Data section (프로그램 초기 데이터 저장)
 - Stack
 - Heap



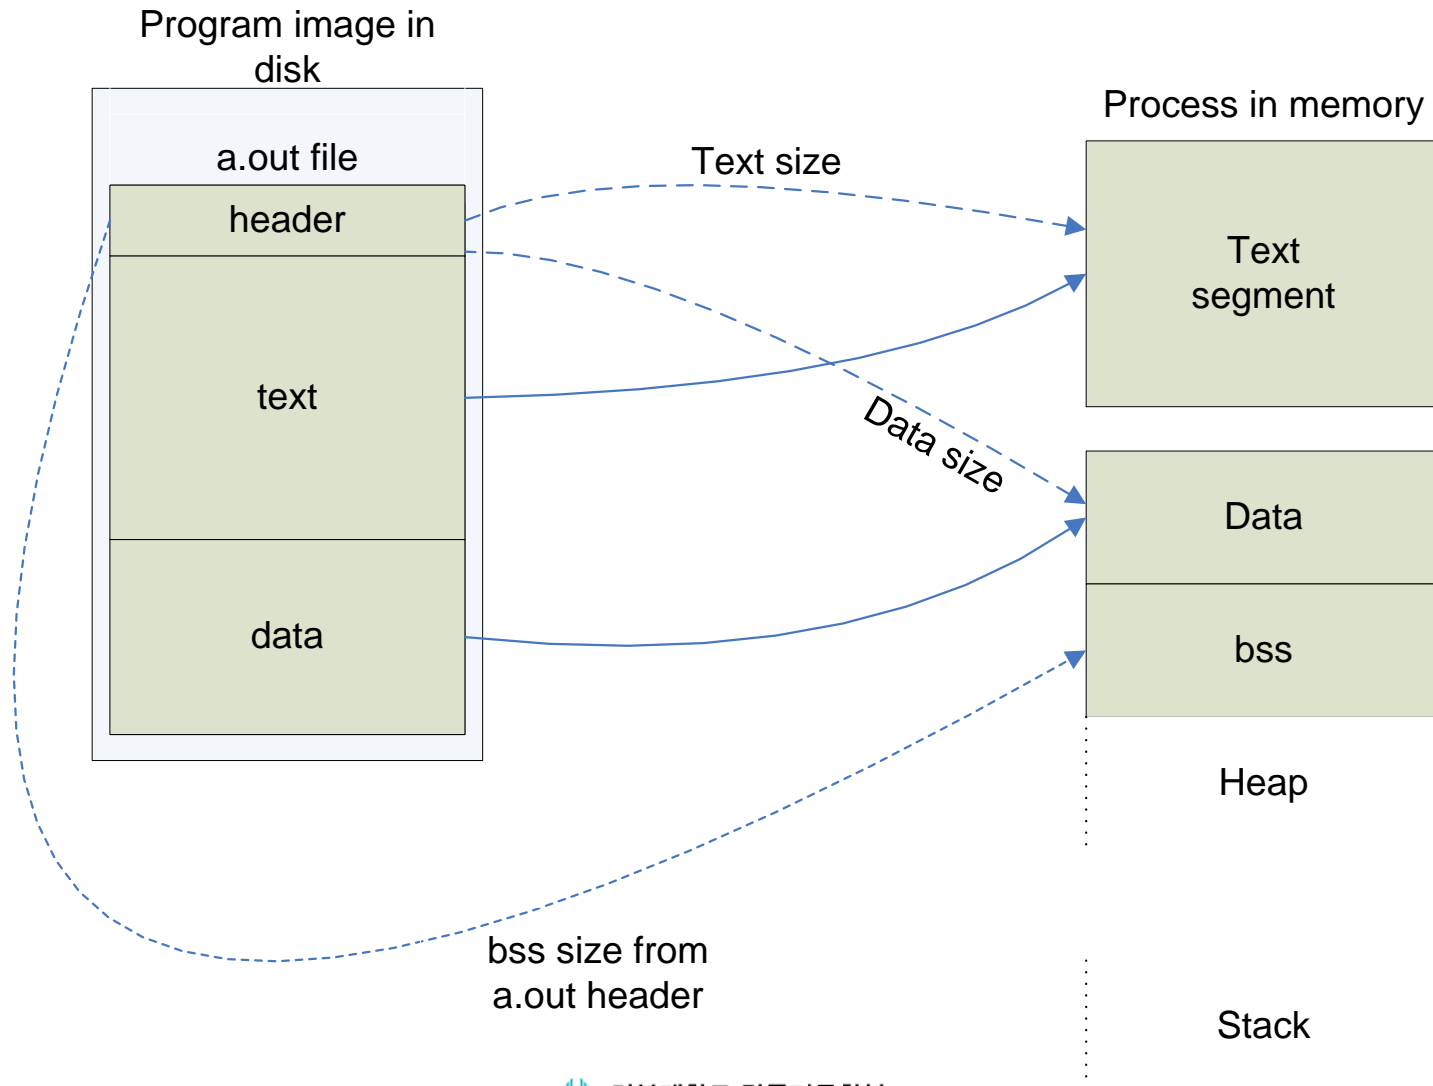
Process Concept

- 유사 용어 정리
 - Job: Batch system
 - Program: 실행되기 전의 프로그램 이미지
 - Passive state (program) vs. Active state (process)
 - 한 프로그램은 여러 프로세스로 구성될 수 있음
 - Task: 일반적 의미로 Job 과 유사.
 - CPU에서의 task: 하나의 실행 흐름. process and thread
 - Thread: Lightweight process. Process보다 가벼운 실행 흐름
 - Processor: CPU 프로세서!

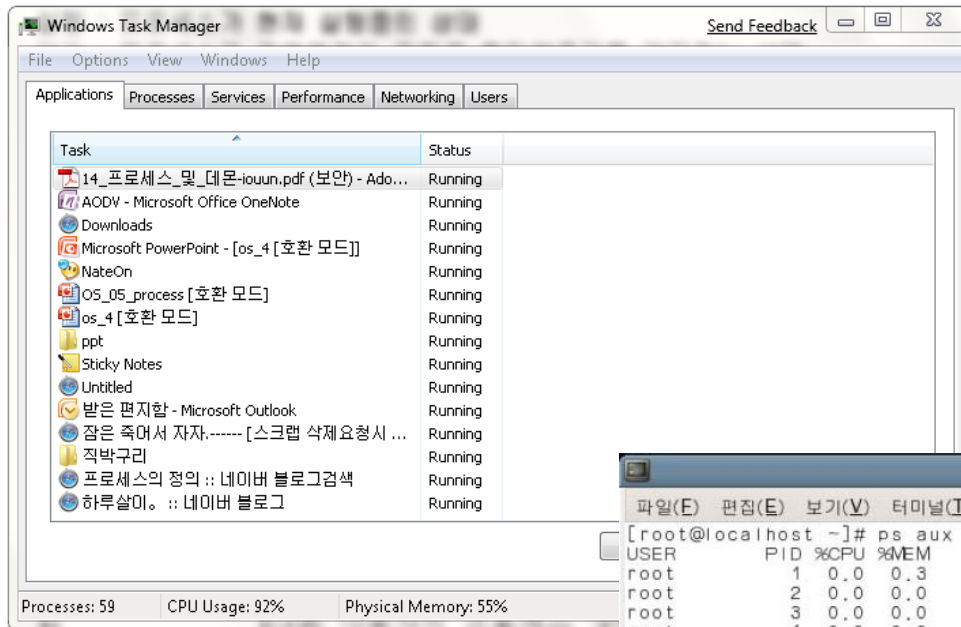
Process in memory



Process from Program

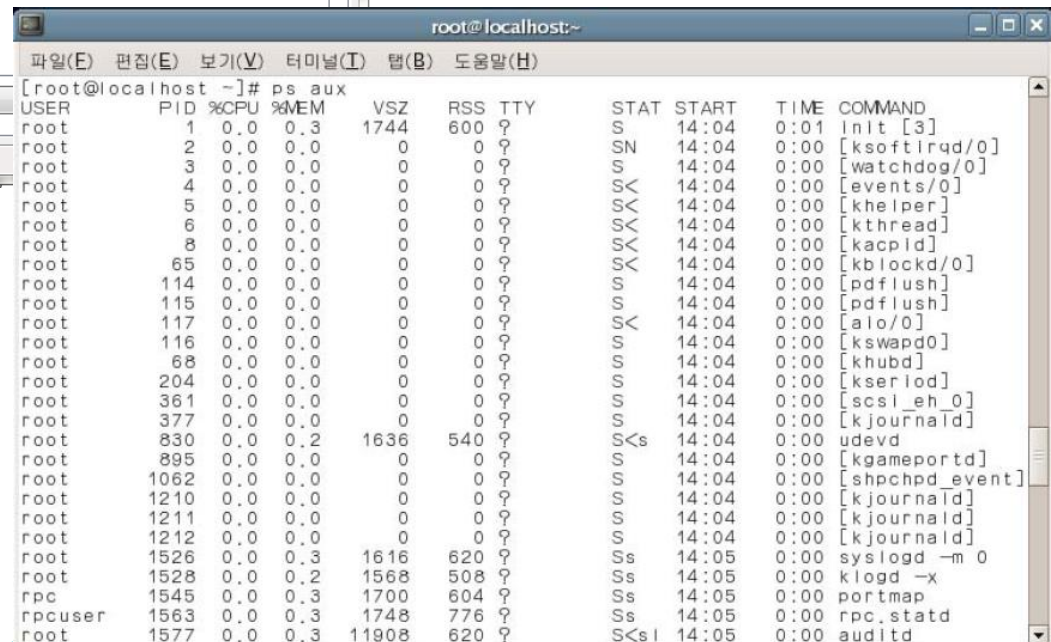


Process from Program



< Windows >

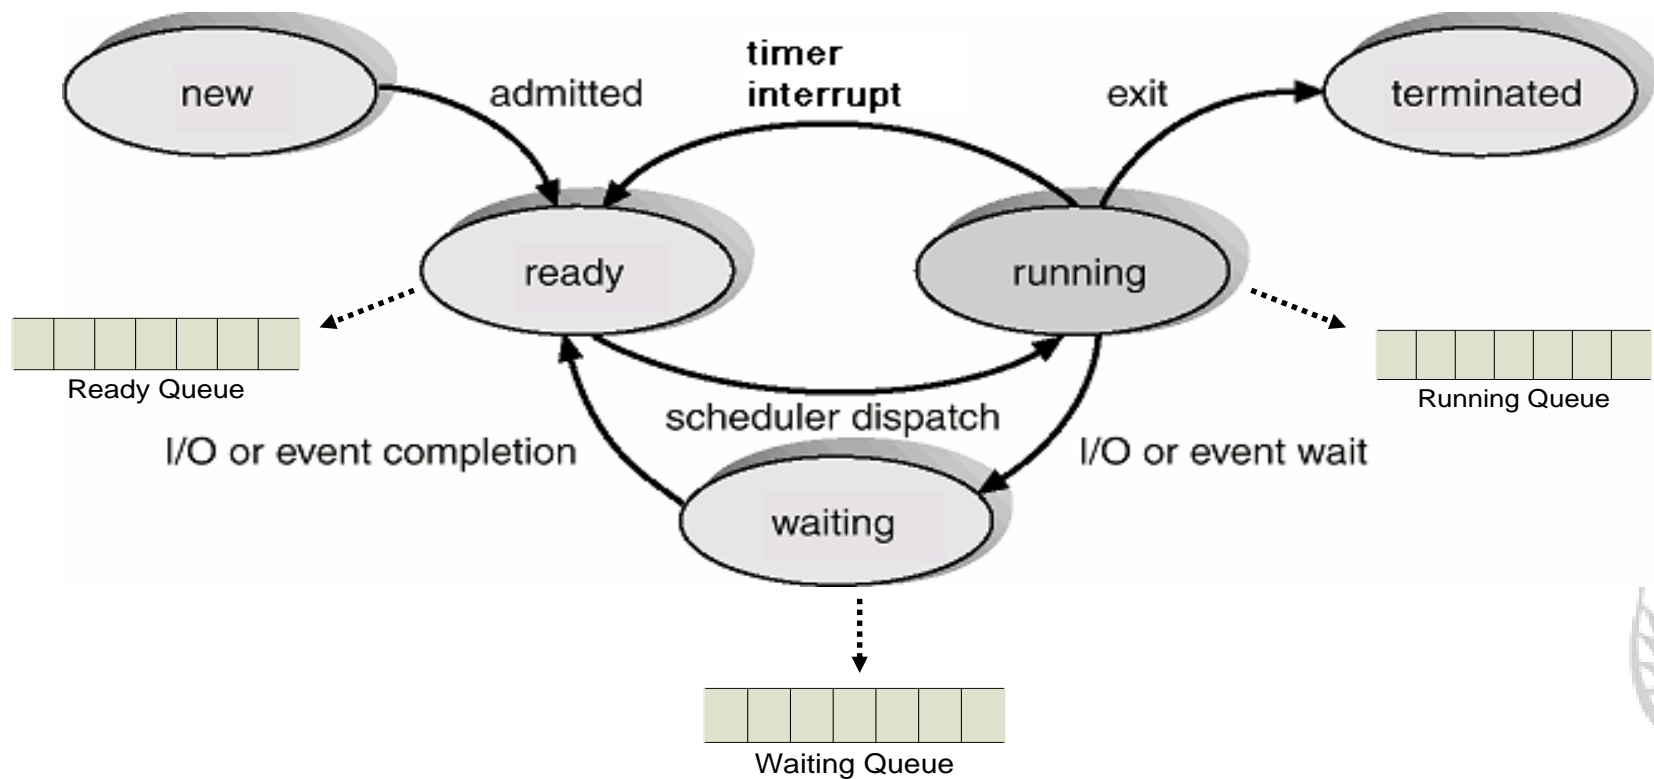
< Linux >



Process State

- New : The process is being created
 - Running : Instruction are being executed
 - Waiting : The process is waiting for some event to occur
 - (such as an I/O completion or reception of a signal)
 - Ready : The process is waiting to be assigned to a processor
 - Terminated: The process has finished execution
-
- 커널 내에 Ready queue, Waiting queue, Running queue를 두고 프로세스들을 상태에 따라 관리한다.

State diagram of process



Process control block

- Process control block (PCB)

- Each process is represented in the OS by PCB.

- Information associated with each process

(also called **task control block**)

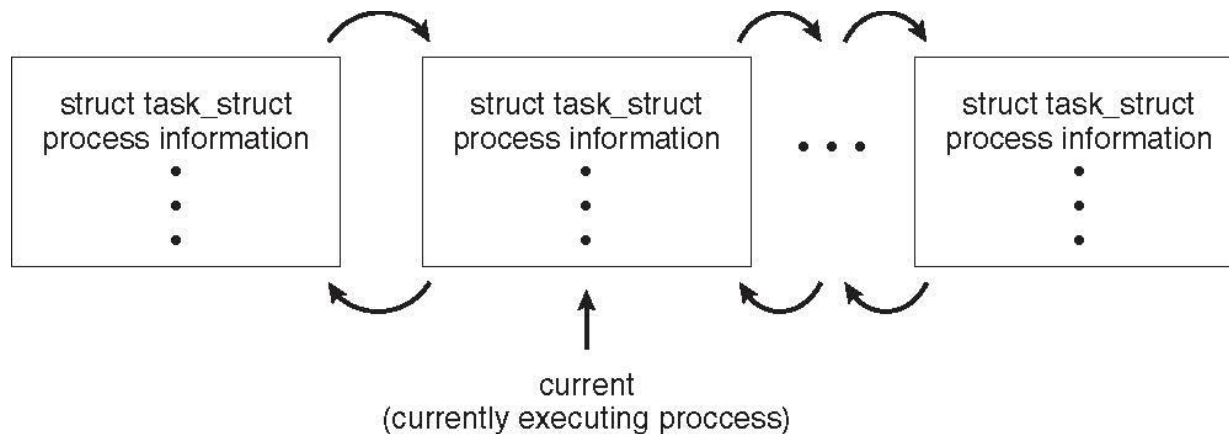
- Process state – running, waiting, etc
 - Program counter – location of instruction to next execute
 - CPU registers – contents of all process-centric registers
 - CPU scheduling information- priorities, scheduling queue pointers
 - Memory-management information – memory allocated to the process
 - Accounting information – CPU used, clock time elapsed since start, time limits
 - I/O status information – I/O devices allocated to process, list of open files
 - And others

process state
process number
program counter
registers
memory limits
list of open files
...



Process Representation in Linux

- Represented by the C structure `task_struct`
- `pid_t pid; /* process identifier */`
`long state; /* state of the process */`
`unsigned int time_slice /* scheduling information */`
`struct task_struct *parent; /* this process's parent */`
`struct list_head children; /* this process's children */`
`struct files_struct *files; /* list of open files */`
`struct mm_struct *mm; /* address space of this process */`



Context Switch (문맥 전환)

- When CPU switches to another process,
 - the system must save the state of the old process
 - and load the saved state for the new process via a **context switch**
- **Context** of a process represented in the PCB
- Context-switch time is overhead; the system does no useful work while switching
 - The more complex the OS and the PCB → the longer the context switch
- Time dependent on hardware support
 - Some hardware provides multiple sets of registers per CPU
→ multiple contexts loaded at once

프로세서 구조에 따른 문맥전환의 차이

- CISC

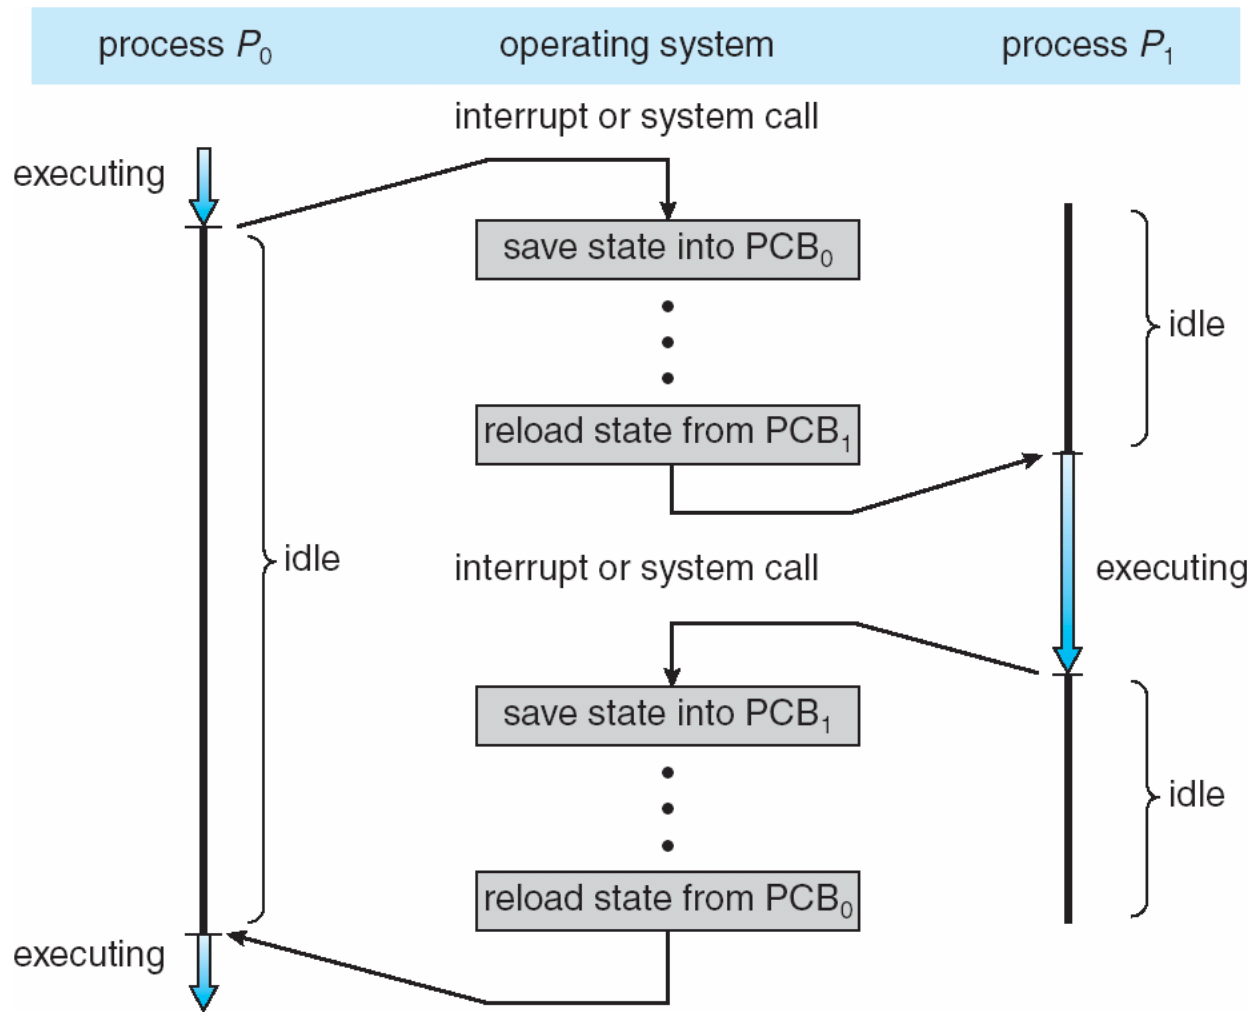
- 복잡한 명령어 셋 구성 -> 효율 높임, 클럭 속도 저하
- 복잡한 회로 -> 물리적인 공간 차지 -> 레지스터 용량 저하
- Ex) Intel Pentium processor

- RISC

- 간단한 명령어 셋 구성 -> 클럭 속도 높임 -> 빠른 수행 속도
- 절약된 물리적 공간에 보다 많은 레지스터 장착
 - 문맥 전환 시 레지스터 내용 변경에 보다 큰 오버헤드가 생김
- 예) Sparc, Arm processor



Context Switch 도식



Operations on Processes

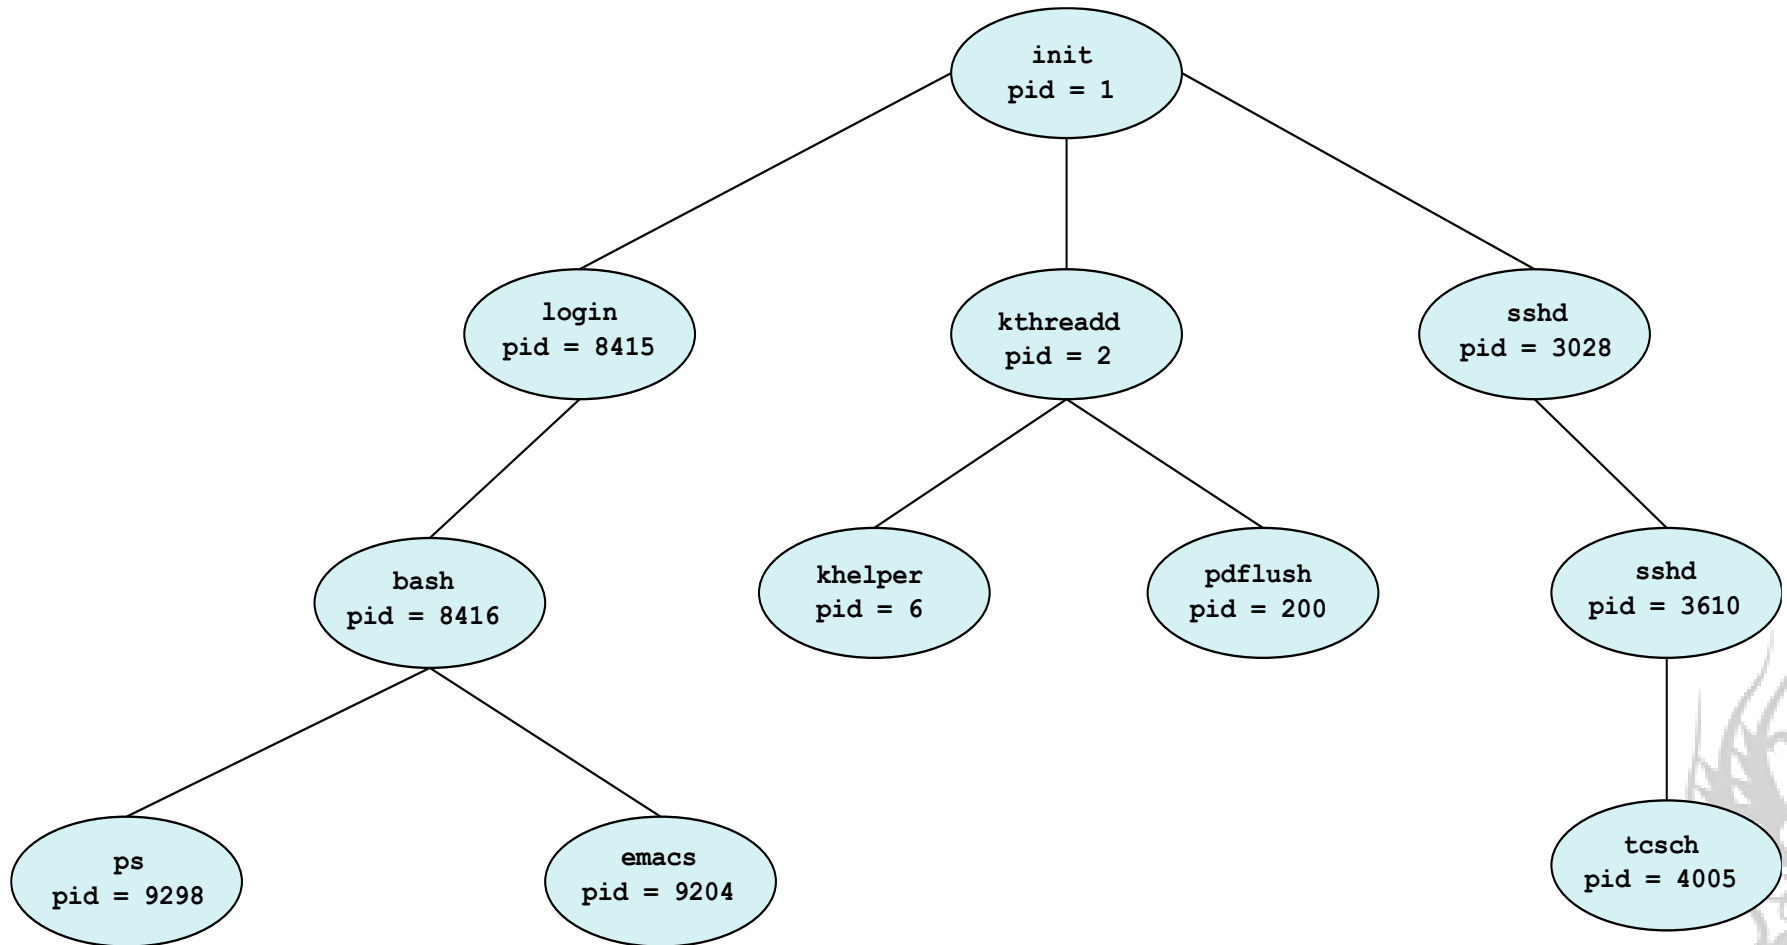
- The processes in the system
 - Can execute concurrently,
 - and they must be created and terminated dynamically.
- System must provide mechanisms for:
 - process creation
 - process termination



Process Creation

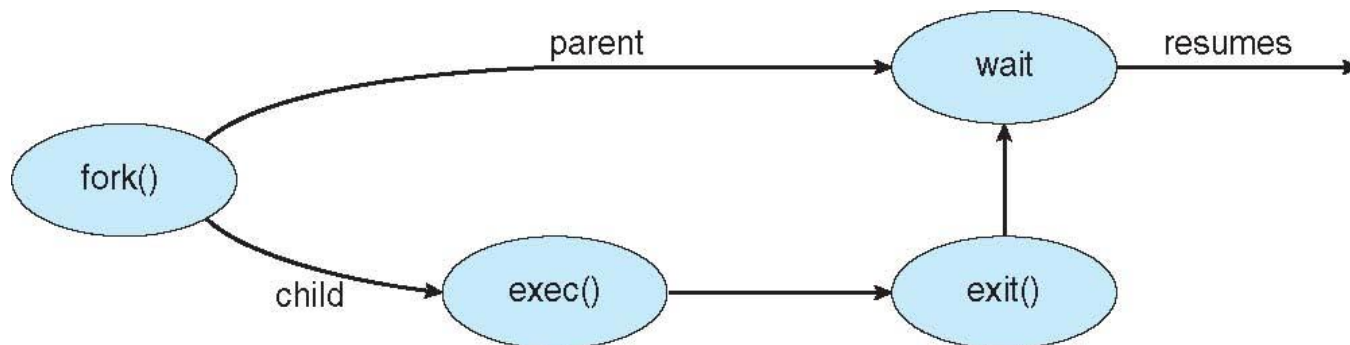
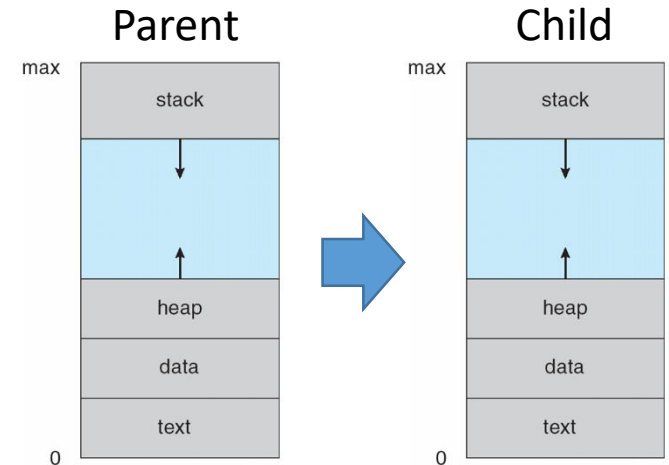
- Parent process create children processes, which, in turn create other processes, forming a tree of processes
- Generally, process identified and managed via a process identifier (pid)
- Resource sharing options
 - Parent and children share all resources
 - Children share subset of parent's resources
 - Parent and child share no resources
- Execution options
 - Parent and children execute concurrently
 - Parent waits until children terminate

A Tree of Processes in Linux

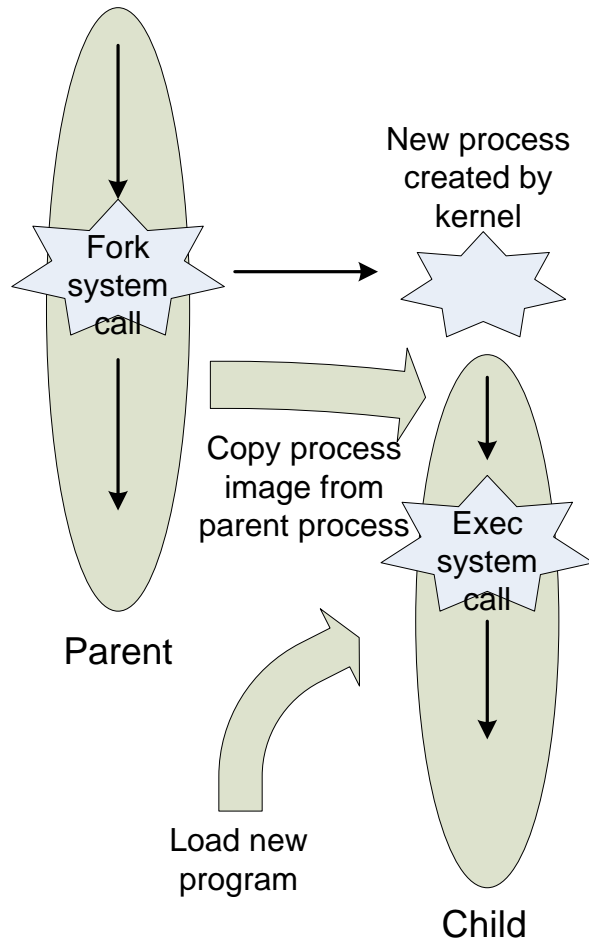


Process Creation (Cont.)

- Address space
 - Child duplicate of parent
 - Child has a program loaded into it
- UNIX examples
 - `fork()` system call creates new process
 - `exec()` system call used after a `fork()` to replace the process' memory space with a new program



Process creation in memory view



- Memory (Text and data section)
- Unix example
 - A new process created by *fork* system call.
 - A new process (child process) consists of a copy of the memory of the original process (parent process)
 - Exec(2): to replace the memory with a new program.

C Program Forking Separate Process

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

int main()
{
    pid_t pid;

    /* fork a child process */
    pid = fork();

    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete");
    }

    return 0;
}
```

Creating a Separate Process via Windows API

```
#include <stdio.h>
#include <windows.h>

int main(VOID)
{
    STARTUPINFO si;
    PROCESS_INFORMATION pi;

    /* allocate memory */
    ZeroMemory(&si, sizeof(si));
    si.cb = sizeof(si);
    ZeroMemory(&pi, sizeof(pi));

    /* create child process */
    if (!CreateProcess(NULL, /* use command line */
        "C:\\WINDOWS\\system32\\mspaint.exe", /* command */
        NULL, /* don't inherit process handle */
        NULL, /* don't inherit thread handle */
        FALSE, /* disable handle inheritance */
        0, /* no creation flags */
        NULL, /* use parent's environment block */
        NULL, /* use parent's existing directory */
        &si,
        &pi))
    {
        fprintf(stderr, "Create Process Failed");
        return -1;
    }
    /* parent will wait for the child to complete */
    WaitForSingleObject(pi.hProcess, INFINITE);
    printf("Child Complete");

    /* close handles */
    CloseHandle(pi.hProcess);
    CloseHandle(pi.hThread);
}
```



Process Termination

- Process executes last statement and then asks the operating system to delete it using the `exit()` system call
 - Returns status data from child to parent (via `wait()`)
 - Process' resources are deallocated by operating system
- Parent may terminate the execution of children processes using the `abort()` system call. Some reasons for doing so:
 - Child has exceeded allocated resources
 - Task assigned to child is no longer required
 - The parent is exiting and the operating systems does not allow a child to continue if its parent terminates

Process Termination

- Some operating systems do not allow child to exist if its parent has terminated. If a process terminates, then all its children must also be terminated.
 - Cascading termination. All children, grandchildren, etc. are terminated.
 - The termination is initiated by the operating system.
- The parent process may wait for termination of a child process by using the `wait()` system call. The call returns status information and the pid of the terminated process
 - `pid = wait(&status);`
- If no parent waiting (did not invoke `wait()`) process is a zombie
- If parent terminated without invoking `wait`, process is an orphan