

# Operating Systems

## 6. CPU: Scheduling (1)

---

Hyunchan, Park

<http://oslab.chonbuk.ac.kr>

Division of Computer Science and Engineering

Chonbuk National University

# Contents

---

- Basic Concepts
  - Process scheduling
  - CPU and I/O bursts
  - Histogram of CPU-burst times
  - Schedulers: short, mid, and long-term
- Scheduling Criteria and Algorithms
  - First-Come, First-Served Scheduling
  - Shortest-Job-First Scheduling
  - Priority Scheduling
  - Round-Robin Scheduling
  - Multilevel Queue Scheduling
  - Multilevel Feedback Queue Scheduling


# Basic Concepts: Process Scheduling

---

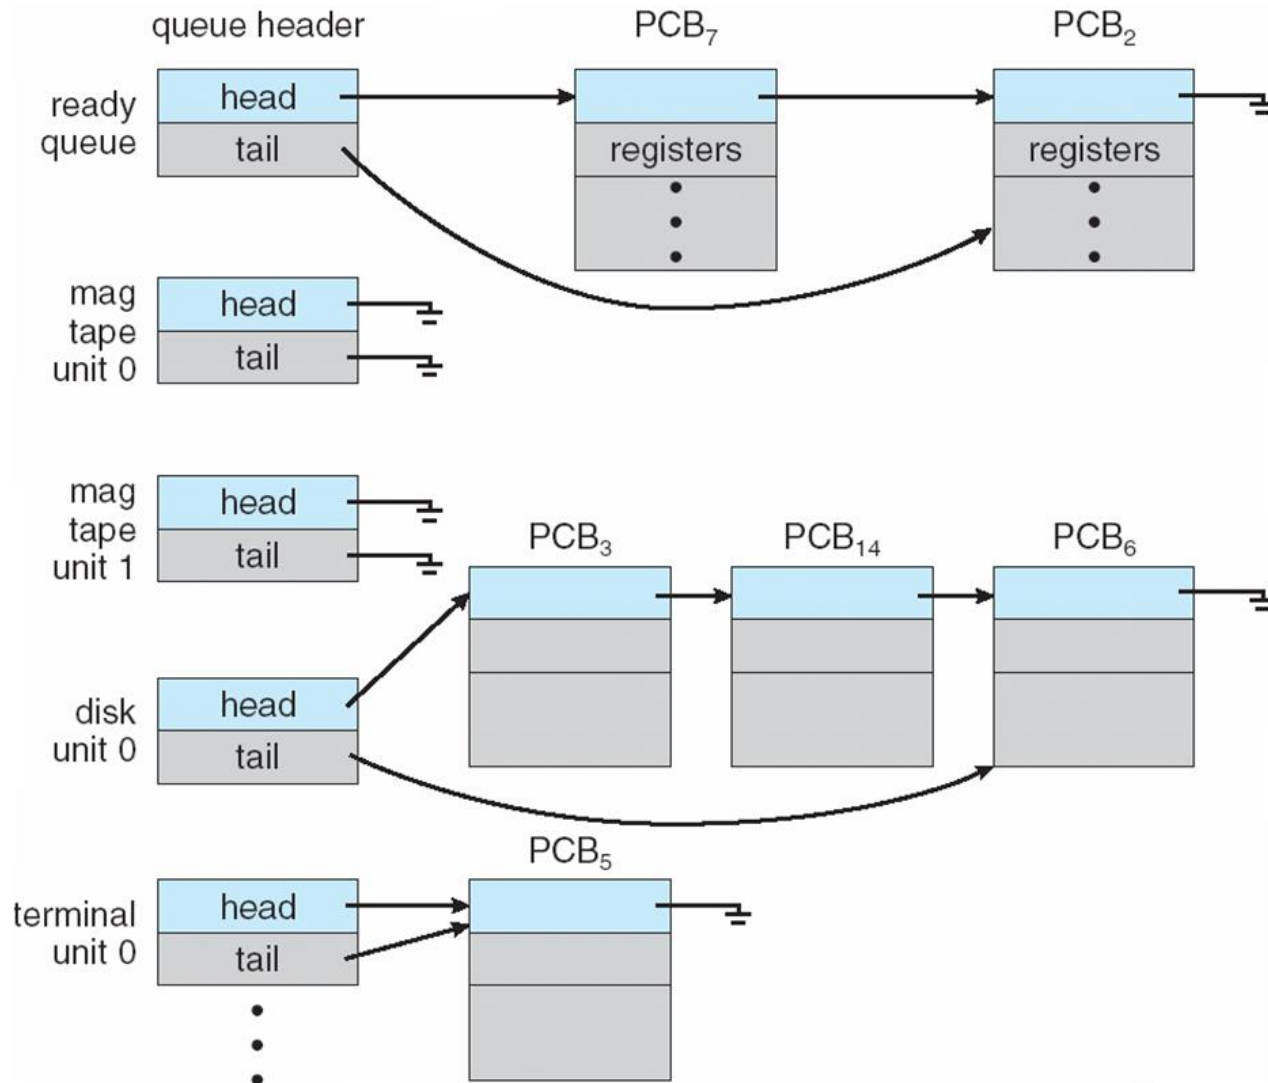
- CPU scheduling: “How to allocate CPU for processes?”
  - To maximize CPU use, quickly switch processes onto CPU for time sharing
- **Process scheduler** selects among available processes for next execution on CPU
- Maintains scheduling queues of processes
  - Job queue – set of all processes in the system
  - Ready queue – set of all processes residing in main memory, ready and waiting to execute
  - Wait queue – set of processes waiting for an I/O device (or device queue)
  - Processes migrate among the various queues

# Representation of Process Scheduling

- Queueing diagram represents queues, resources, flows

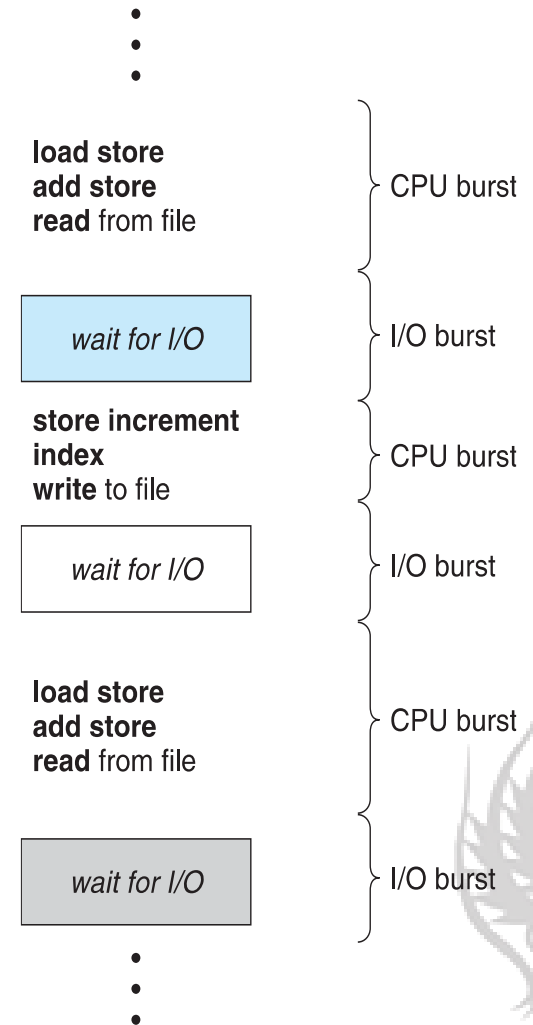


# Ready Queue And Various I/O Device Queues




# Basic Concepts: Process Scheduling

- Maximum CPU utilization obtained with multiprogramming
- CPU–I/O Burst Cycle: Process execution consists of a **cycle** of CPU execution and I/O wait
- **CPU burst** followed by **I/O burst**
  - CPU burst distribution is of main concern
- Processes can be described as either:
  - I/O-bound process – spends more time doing I/O than computations, many short CPU bursts
  - CPU-bound process – spends more time doing computations; few very long CPU bursts



# Histogram of CPU-burst Times

- Common characteristics
  - A large number of short CPU bursts
  - A small number of long CPU bursts



# Schedulers

---

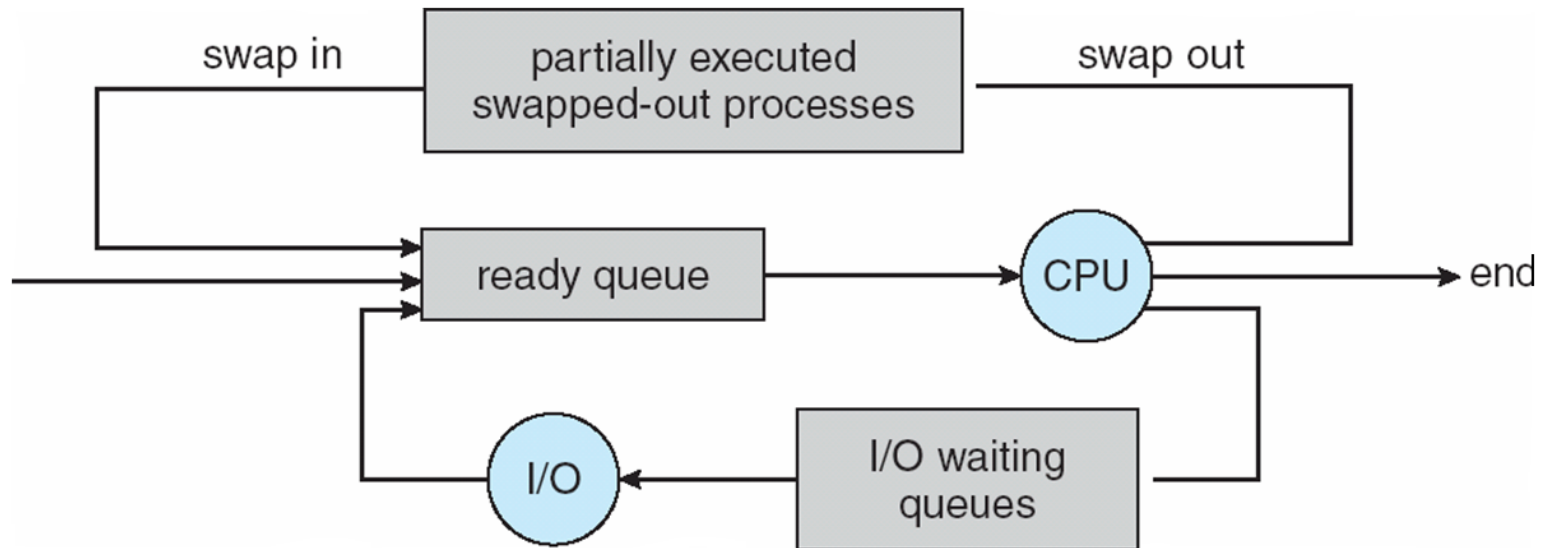
- Short-term scheduler (or CPU scheduler) – selects which process should be executed next and allocates CPU
  - Sometimes the only scheduler in a system
  - Short-term scheduler is invoked frequently (milliseconds)  $\Rightarrow$  (must be fast)
- Long-term scheduler (or job scheduler) – selects which processes should be brought into the ready queue
  - Long-term scheduler is invoked infrequently (seconds, minutes)  $\Rightarrow$  (may be slow)
  - The long-term scheduler controls the degree of multiprogramming
- Long-term scheduler strives for good process mix
  - Among CPU and I/O-bound processes





# Addition of Medium Term Scheduling

- Medium-term scheduler can be added if degree of multiple programming needs to decrease
  - In other words, when the memory is not enough
  - Remove process from memory, store on disk, bring back in from disk to continue execution: swapping



# Short-term Scheduler (CPU Scheduler)

---

- Short-term scheduler selects from among the processes in ready queue, and allocates the CPU to one of them
  - Queue may be ordered in various ways
- CPU scheduling decisions may take place when a process:
  - 1. Switches from running to waiting state
  - 2. Switches from running to ready state
  - 3. Switches from waiting to ready
  - 4. Terminates

# Non-preemptive and preemptive schedulers

---

- Non-preemptive (비선점형)
  - Scheduling under 1 and 4
  - OS cannot control the executing process
- All other scheduling is preemptive (선점형)
  - OS can control or interrupt the executing process
  - Issue: synchronization
    - Consider access to shared data
    - Consider preemption while in kernel mode
    - Consider interrupts occurring during crucial OS activities

# Dispatcher

---

- Dispatcher module gives control of the CPU to the process selected by the short-term scheduler; this involves:
  - switching context
  - switching to user mode
  - jumping to the proper location in the user program to restart that program
- **Dispatch latency** – time it takes for the dispatcher to stop one process and start another running
  - Context switching overhead

---

# Criteria and Algorithms



# Scheduling Criteria

---

- **CPU utilization** – keep the CPU as busy as possible
- **Throughput** – # of processes that complete their execution per time unit
- **Turnaround time** – amount of time to execute a particular process (waiting time is included)
- **Waiting time** – amount of time a process has been waiting in the ready queue
- **Response time** – amount of time it takes from when a request was submitted until the first response is produced, not output (for time-sharing environment)

# Scheduling Algorithm Optimization Criteria

---

- Max CPU utilization
- Max throughput
- Min turnaround time
- Min waiting time
- Min response time
- It's impossible to satisfy all the criteria in the singular scheduler
- Only focuses on the main requirement depend on the system
  - Super computer: CPU utilization
  - Main frame or work station: Throughput, turnaround time
  - Personal computer: Response time



# Scheduling Algorithms

---

- First-Come, First-Served Scheduling
- Shortest-Job-First Scheduling
- Priority Scheduling
- Round-Robin Scheduling
- Multilevel Queue Scheduling
- Multilevel Feedback Queue Scheduling



# First- Come, First-Served (FCFS) Scheduling

<u>Process</u>	<u>Burst Time</u>
$P_1$	24
$P_2$	3
$P_3$	3

- Suppose that the processes arrive in the order:  $P_1, P_2, P_3$   
The Gantt Chart for the schedule is:



- Waiting time for  $P_1 = 0$ ;  $P_2 = 24$ ;  $P_3 = 27$
- Average waiting time:  $(0 + 24 + 27)/3 = 17$

# FCFS Scheduling (Cont.)

Suppose that the processes arrive in the order:

$$P_2, P_3, P_1$$

- The Gantt chart for the schedule is:



- Waiting time for  $P_1 = 6$ ;  $P_2 = 0$ ;  $P_3 = 3$
- Average waiting time:  $(6 + 0 + 3)/3 = 3$
- Much better than previous case
- Convoy effect** - short process behind long process
  - Consider one CPU-bound and many I/O-bound processes

# Shortest-Job-First (SJF) Scheduling

---

- Associate with each process the length of its next CPU burst
  - Use these lengths to schedule the process with the shortest time
- SJF is optimal – gives minimum average waiting time for a given set of processes
  - The difficulty is knowing the length of the next CPU request
  - Could ask the user



# Example of SJF

Process

Burst Time

$P_1$

6

$P_2$

8


$P_3$

7

$P_4$

3

- SJF scheduling chart



- Average waiting time =  $(3 + 16 + 9 + 0) / 4 = 7$




# Example of Shortest-remaining-time-first

- Now we add the concepts of varying arrival times and preemption to the analysis

Process	Arrival Time	Burst Time
P1	0	8
P2	1	4
P3	2	9
P4	3	5

- Preemptive SJF Gantt Chart



- Average waiting time =  $[(10-1)+(1-1)+(17-2)+5-3]/4 = 26/4 = 6.5$  msec

# Priority Scheduling

---

- A priority number (integer) is associated with each process
- The CPU is allocated to the process with the highest priority (smallest integer  $\equiv$  highest priority)
- SJF is priority scheduling where priority is the inverse of CPU burst time
- Problem  $\equiv$  Starvation – low priority processes may never execute
- Solution  $\equiv$  Aging – as time progresses increase the priority of the process

# Example of Priority Scheduling

<u>Process</u>	<u>Burst Time</u>	<u>Priority</u>
$P_1$	10	3
$P_2$	1	1
$P_3$	2	4
$P_4$	1	5
$P_5$	5	2

- Priority scheduling Gantt Chart



- Average waiting time = 8.2 msec
  - $= (0+1+6+16+18)/5$

# Round Robin (RR)

---


- Each process gets a small unit of CPU time, usually 1-100 milliseconds. After this time has elapsed, the process is preempted and added to the end of the ready queue
  - Time quantum: the smallest unit of allocation (e.g. 1 ms)
    - H/W timer interrupts every quantum to enter the kernel mode, and kernel schedules next process if it is necessary
  - Time slice: allocated (or allowed to use CPU) time for a process in a round (e.g. TS = 10 ms = 10 of time quantum)
- Performance
  - *Time slice* large  $\Rightarrow$  FIFO (high performance due to minimal context switches)
  - *Time slice* small  $\Rightarrow$  High responsiveness but low performance
  - *Time slice* must be large with respect to context switch, otherwise overhead is too high



# Example of RR with Time Slice = 4


<u>Process</u>	<u>Burst Time</u>
$P_1$	24
$P_2$	3
$P_3$	3

- The Gantt chart is:



- Typically, higher average turnaround, but better **response**
- Time slice should be large compared to context switch time
- Time slice usually 1ms to 10ms, context switch < 10 usec

# Time slice and Context Switch Time



# Multilevel Queue

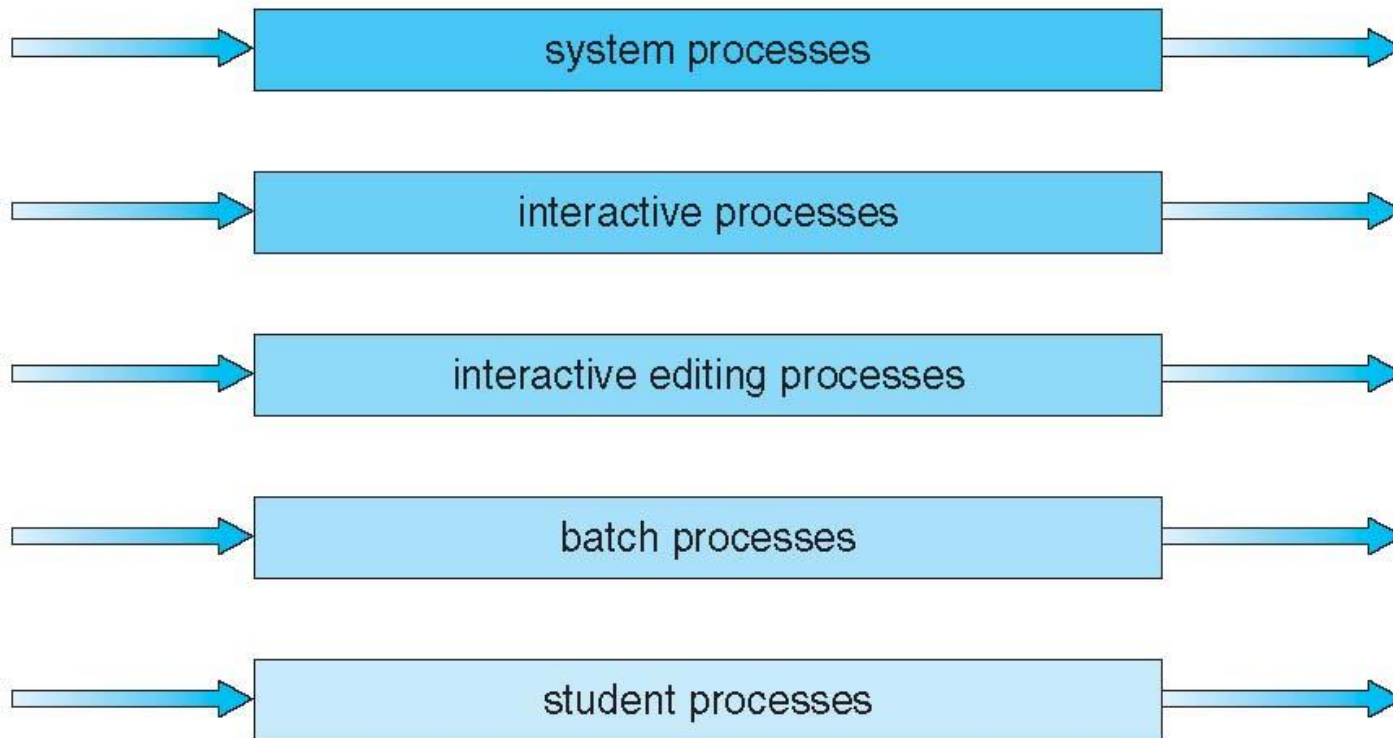
---

- Ready queue is partitioned into separate queues, eg:
  - **foreground** (interactive) w/ RR
  - **background** (batch) w/FCFS
- Process permanently in a given queue
- Each queue has its own scheduling algorithm
- Scheduling must be done between the queues:
  - Fixed priority scheduling; (i.e., serve all from foreground then from background). Possibility of starvation.
  - Time slice – each queue gets a certain amount of CPU time which it can schedule amongst its processes; i.e., 80% to foreground in RR
  - 20% to background in FCFS

# Multilevel Queue Scheduling

---

highest priority



lowest priority



# Multilevel Feedback Queue


---

- A process can move between the various queues; aging can be implemented this way
- Multilevel-feedback-queue scheduler defined by the following parameters:
  - number of queues
  - scheduling algorithms for each queue
  - method used to determine when to upgrade a process
  - method used to determine when to demote a process
  - method used to determine which queue a process will enter when that process needs service

# Example of Multilevel Feedback Queue

- Three queues:

- Q0 – RR with time slice 8 milliseconds
- Q1 – RR time slice 16 milliseconds
- Q2 – FCFS



- Scheduling

- A new job enters queue Q0 which is served FCFS
  - When it gains CPU, job receives 8 milliseconds
  - If it does not finish in 8 milliseconds, job is moved to queue Q1
- At Q1 job is again served FCFS and receives 16 additional milliseconds
  - If it still does not complete, it is preempted and moved to queue Q2