

Operating Systems

8. CPU: Thread

Hyunchan, Park

<http://oslab.chonbuk.ac.kr>

Division of Computer Science and Engineering

Chonbuk National University

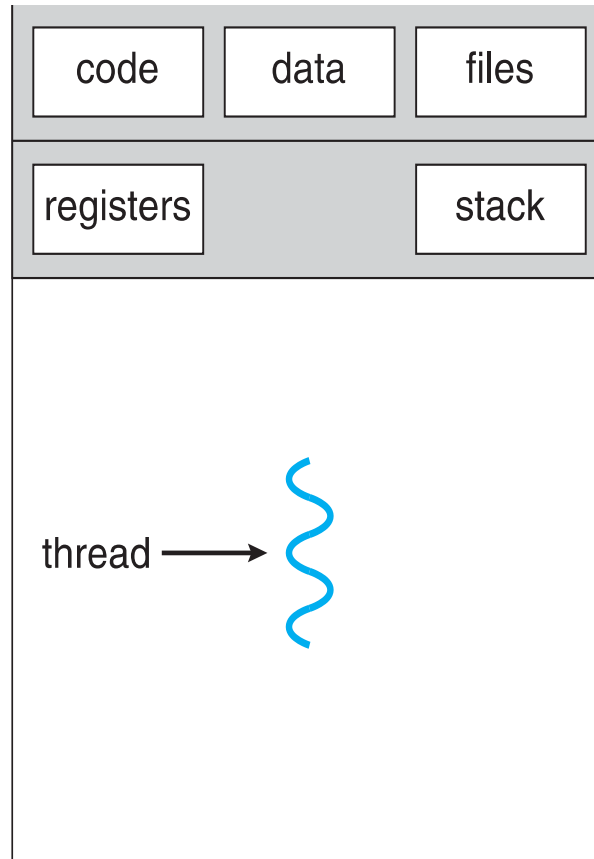
Contents

- Thread
- User and Kernel threads: Multithreading Models
 - Many-to-One
 - One-to-One
 - Many-to-Many
- Issues with threads
 - Creation (fork and exec system calls)
 - Cancellation
 - Thread pools
 - IPC between threads
 - Signal handling
- Implementation: Linux thread

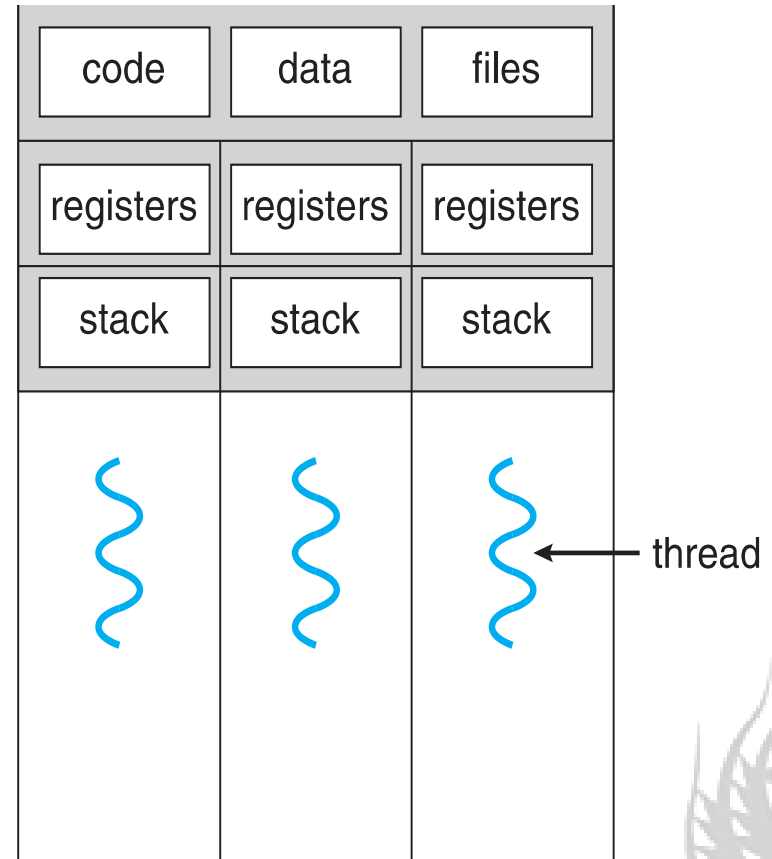
Thread

- So far, process has a single thread of execution
- Consider having multiple program counters per process
 - Multiple locations can execute at once
 - Multiple threads of control -> threads
 - Must then have storage for thread details, multiple program counters in PCB
- Thread
 - The execution unit in a process
 - Program counter, register set, stack
 - Code, data, and opened files are shared among threads in a process

Single and Multithreaded Processes



single-threaded process



multithreaded process



Processes vs. Threads

- Processes

- Protection domain between processes
 - Cannot directly access the other's memory \Rightarrow must use IPC
- Heavy weight: more operations are needed for context switch


- Threads

- Code and data sections are shared among threads in a process
- Light weight: more efficient switching between threads in a process
 - Because they share the same address space
 - No TLB and cache flush (will be explained later)

Motivation

- Most modern applications are multithreaded
- Threads run within application
- Multiple tasks with the application can be implemented by separate threads
 - Update display
 - Fetch data
 - Spell checking
 - Answer a network request
- Process creation is heavy-weight while thread creation is light-weight
- Can simplify code, increase efficiency
- Kernels are generally multithreaded

Multithreaded Server Architecture



Benefits

- Responsiveness
 - may allow continued execution if part of process is blocked, especially important for user interfaces
- Resource Sharing
 - threads share resources of process, easier than shared memory or message passing
- Performance
 - cheaper than process creation, thread switching lower overhead than context switching
- Scalability
 - process can take advantage of multiprocessor architectures

Amdahl's Law

- Identifies performance gains from adding additional cores to an application that has both serial and parallel components
- S is serial portion
- N processing cores

$$\text{speedup} \leq \frac{1}{S + \frac{(1-S)}{N}}$$

- That is, if application is 75% parallel / 25% serial, moving from 1 to 2 cores results in speedup of 1.6 times
- As N approaches infinity, speedup approaches $1 / S$
- Serial portion of an application has disproportionate effect on performance gained by adding additional cores
- But does the law take into account contemporary multicore systems?

User Threads and Kernel Threads

- User threads
 - Management done by **user-level threads library**
 - Three primary thread libraries:
 - POSIX Pthreads, Windows threads, Java threads
- Pros
 - High performance: only maintains user-mode context
- Cons
 - If a thread executes a system call, entire threads are blocked by kernel
 - Because they are just a sole process for the kernel: kernel doesn't know about user level threads

User Threads and Kernel Threads

- Kernel threads
 - Management done by **the Kernel**
 - Examples – virtually all general purpose operating systems, including:
 - Windows, Solaris, Linux, Tru64 UNIX, and Mac OS X
- Pros
 - Support high level of parallelism
 - Each of threads use a system call exclusively; does not blocked by other thread's system call
- Cons
 - High cost on thread creation/management
 - Should maintains kernel level context and metadata
 - E.g. have to enlarge the PCB for the threads


User to Kernel: Multithreading Models

- Many-to-One
- One-to-One
- Many-to-Many



Many-to-One

- Many user-level threads mapped to single kernel thread
- Cons: One thread blocking causes all to block
 - Multiple threads may not run in parallel on muticore system because only one may be in kernel at a time
- Few systems currently use this model
 - Who does not support kernel threads
 - Examples:
 - Solaris Green Threads
 - GNU Portable Threads




One-to-One

- Each user-level thread maps to kernel thread
- Creating a user-level thread creates a kernel thread
- Pros: More concurrency than many-to-one
- Cons: Number of threads per process sometimes restricted due to overhead


- Examples

- Windows
- Linux
- Solaris 9 and later




Many-to-Many Model

- Allows many user level threads to be mapped to many kernel threads
- Pros: Allows the operating system to create a sufficient number of kernel threads
- Solaris prior to version 9
- Windows with the *ThreadFiber* package



Two-level Model

- Similar to M:M, except that it allows a user thread to be bound to kernel thread
- Examples
 - IRIX
 - HP-UX
 - Tru64 UNIX
 - Solaris 8 and earlier



Threading Issues

- Semantics of fork() and exec() system calls
- Thread cancellation of target thread
 - Asynchronous or deferred
- Thread pools
- IPC between threads
- Signal handling
 - Synchronous and asynchronous



Semantics of fork() and exec()

- Does fork() duplicate only the calling thread or all threads?
 - Some UNIXes have two versions of fork
- exec() usually works as normal
 - Replace the running process including all threads
- Is it a right policy?
 - How about that the only thread who execute fork() should be replaced by exec()?

Thread Cancellation

- Terminating a thread before it has finished
 - Thread to be canceled is target thread
- Two general approaches:
 - Asynchronous cancellation terminates the target thread immediately
 - Deferred cancellation allows the target thread to periodically check if it should be cancelled
- Pthread code to create and cancel a thread:

```
pthread_t tid;

/* create the thread */
pthread_create(&tid, 0, worker, NULL);

. . .

/* cancel the thread */
pthread_cancel(tid);
```

Thread Cancellation (Cont.)

- Invoking thread cancellation requests cancellation, but actual cancellation depends on thread state

Mode	State	Type
Off	Disabled	–
Deferred	Enabled	Deferred
Asynchronous	Enabled	Asynchronous

- If thread has cancellation disabled, cancellation remains pending until thread enables it
- Default type is deferred
 - Cancellation only occurs when thread reaches cancellation point
 - I.e. `pthread_testcancel()`
 - Then cleanup handler is invoked
- On Linux systems, thread cancellation is handled through signals

Thread Pools

- Create a number of threads in a pool where they await work
- Advantages:
 - Usually slightly faster to service a request with an existing thread than create a new thread
 - Allows the number of threads in the application(s) to be bound to the size of the pool
 - Separating task to be performed from mechanics of creating task allows different strategies for running task
 - i.e. Tasks could be scheduled to run periodically

IPC between threads

- Shared memory is most appropriate
 - Because threads share a same address space
 - High performance, but synchronization is required
- Threads naturally decrease the requirements for IPCs
- The IPC between threads in different processes?
 - No differences with IPC between processes
 - Originated in POOR program design

Signal Handling

- Signals are used in UNIX systems to notify a process that a particular event has occurred.
- A signal handler is used to process signals
 - Signal is generated by particular event
 - Signal is delivered to a process
 - Signal is handled by one of two signal handlers:
 - default
 - user-defined
- Every signal has default handler that kernel runs when handling signal
 - User-defined signal handler can override default
 - For single-threaded, signal delivered to process

Signal Handling (Cont.)

- Where should a signal be delivered for multi-threaded?
 - Deliver the signal to the thread to which the signal applies
 - Deliver the signal to every thread in the process
 - Deliver the signal to certain threads in the process
 - Assign a specific thread to receive all signals for the process



Implementation

- Thread library
- POSIX Pthread
 - IEEE 1003.1c: pthread_create()
- Windows threads API
 - Via system call: CreateThread()
- Linux threads
 - Introduced in version 2.2
 - Via system call: clone()
- Java threads
 - Thread class

Linux Threads

- Linux refers to them as tasks rather than threads
- Thread creation is done through clone() system call
- clone() allows a child task to share the address space of the parent task (process)
 - Flags control behavior

flag	meaning
CLONE_FS	File-system information is shared.
CLONE_VM	The same memory space is shared.
CLONE_SIGHAND	Signal handlers are shared.
CLONE_FILES	The set of open files is shared.

- struct task_struct points to process data structures (shared or unique)

Multicore Programming


- Multicore or multiprocessor systems putting pressure on programmers, challenges include:
 - Dividing activities
 - Balance
 - Data splitting
 - Data dependency
 - Testing and debugging
- **Parallelism** implies a system can perform more than one task simultaneously
- **Concurrency** supports more than one task making progress
 - Single processor / core, scheduler providing concurrency

Multicore Programming (Cont.)

- Types of parallelism
 - Data parallelism – distributes subsets of the same data across multiple cores, same operation on each
 - Task parallelism – distributing threads across cores, each thread performing unique operation
- As # of threads grows, so does architectural support for threading
 - CPUs have cores as well as hardware threads
 - Consider Oracle SPARC T4 with 8 cores, and 8 hardware threads per core

Concurrency vs. Parallelism

- Concurrent execution on single-core system:



- Parallelism on a multi-core system:

