

Operating Systems

11. Memory: Main Memory

Hyunchan, Park

<http://oslab.chonbuk.ac.kr>

Division of Computer Science and Engineering

Chonbuk National University

Contents

- Background
 - Address binding Schemes: Compile, Load and Execution time
 - Logical and physical memory
 - Address Translation w/ MMU: Dynamic relocation
- Memory Allocation: Contiguous and Multiple partition
 - Issues: Swapping and Fragmentation

Main memory




DDR3



DDR4



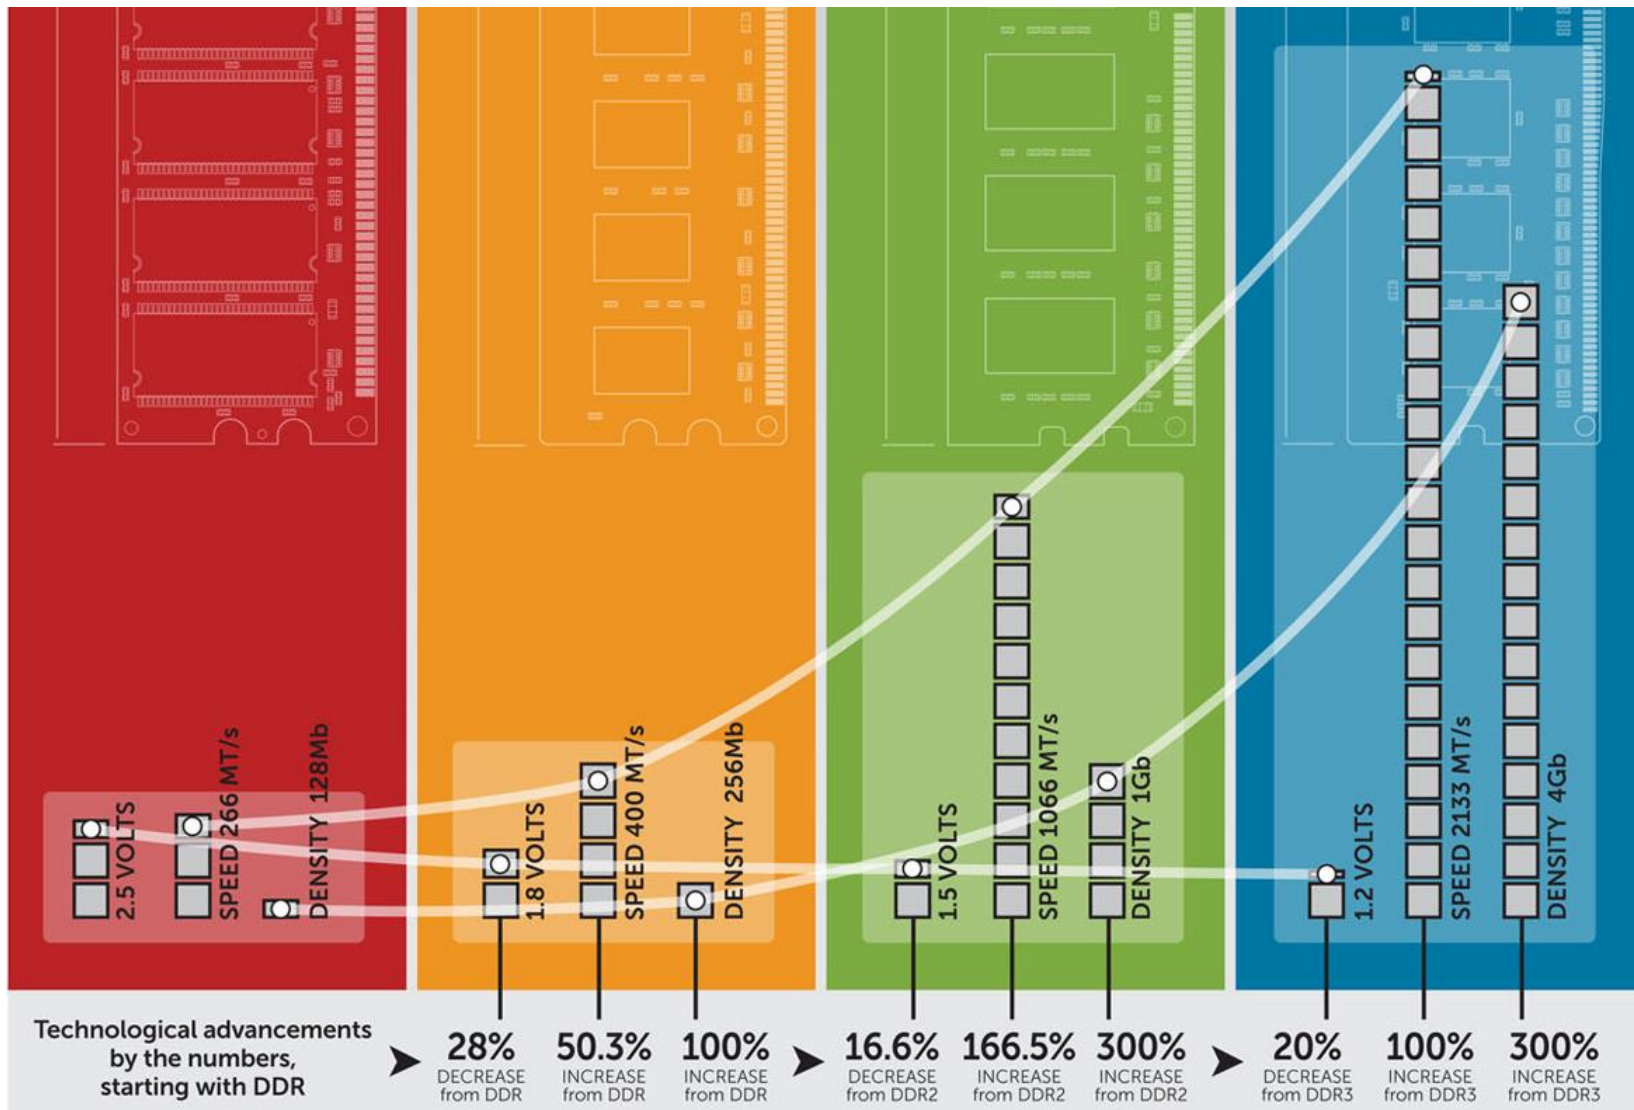
DDR4
HIGH-PERFORMANCE MEMORY



전북대학교 컴퓨터공학부

Division of Computer Science and Engineering
Chonbuk National University

Main memory



Background


- Program must be brought (from) into and placed within a process for it to be run
- M and are only storage CPU can access directly
- Memory unit only sees a stream of
 - addresses + read requests, or
 - address + data and write requests
- Register access in CPU clock (or less)
- Access for can take many cycles, causing a stall
- C sits between main memory and CPU registers
- Protection of memory required to ensure correct operation

Background

- Program must be brought (from **disk**) into **memory** and placed within a process for it to be run
- **Main memory** and **registers** are only storage CPU can access directly
- Memory unit only sees a stream of
 - addresses + read requests, or
 - address + data and write requests
- Register access in **one** CPU clock (or less)
- Access for **main memory** can take many cycles, causing a stall
- **Cache** sits between main memory and CPU registers
- Protection of memory required to ensure correct operation

Remind: Batch and Multiprogramming

- 2개 이상의 작업을 동시에 실행
 - 운영체제는 여러 개의 작업을 메모리에 동시에 유지
 - 현재 실행중인 작업이 I/O를 할 경우 다음 작업을 순차적 실행
 - 스케줄링 고려사항 – first-come, first-served
- 장점: CPU 활용률 극대화
 - CPU가 쉬는 시간이 없음

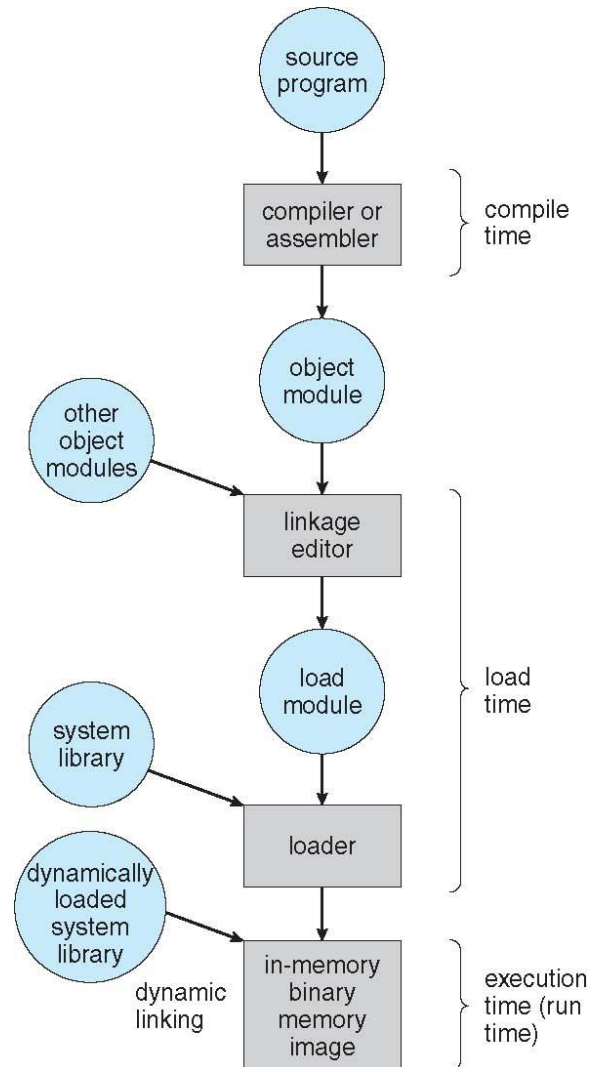


Address Binding

- Programs on disk, ready to be brought into memory to execute form an input queue
 - Without support, must be loaded into address 0x0000
- Inconvenient to have first user process physical address always at 0x0000
 - How can it not be?
- Further, addresses represented in different ways at different stages of a program's life
 - Source code addresses usually symbolic
 - Compiled code addresses bind to relocatable addresses
 - i.e. “14 bytes from beginning of this module”
 - Linker or loader will bind relocatable addresses to absolute addresses
 - i.e. 74014
 - Each binding maps one address space to another



Multistep Processing of a User Program



FYI. Dynamic Linking

- Static linking – system libraries and program code combined by the loader into the binary program image
- Dynamic linking – linking postponed until execution time
- Small piece of code, stub, used to locate the appropriate memory-resident library routine
- Stub replaces itself with the address of the routine, and executes the routine
- Operating system checks if routine is in processes' memory address
 - If not in address space, add to address space
- Dynamic linking is particularly useful for libraries
 - System also known as shared libraries

Binding of Instructions and Data to Memory

- Address binding of instructions and data to memory addresses can happen at three different stages
 - Compile time: If memory location known a priori, absolute code can be generated; must recompile code if starting location changes
 - Load time: Must generate relocatable code if memory location is not known at compile time
 - Execution time: Binding delayed until run time if the process can be moved during its execution from one memory segment to another
 - Need hardware support for address maps (e.g., base and limit registers)

Logical vs. Physical Address Space

- The concept of a logical address space that is bound to a separate physical address space is central to proper memory management
 - Logical address – generated by the CPU; also referred to as virtual address
 - The set of all logical addresses generated by a program
 - Physical address – address seen by the memory unit
- Logical and physical addresses are the same
 - in compile-time and load-time address-binding schemes;
- Logical (virtual) and physical addresses differ
 - in execution-time address-binding scheme


Virtual address space

- Virtual address space – logical view of how process is stored in memory
 - Usually start at address 0, contiguous addresses until end of space
 - MMU must map logical to physical
- The user program deals with logical addresses; it never sees the real physical addresses
 - Execution-time binding occurs when reference is made to location in memory
 - Logical address bound to physical addresses
- Virtual memory can be implemented via:
 - Demand paging and Demand segmentation (will be covered next)

Address translation with Memory-Management Unit (MMU)

- MMU: Hardware device that at run time maps virtual to physical address
 - Many methods possible, covered in the rest of this chapter
- To start, consider simple scheme where the value in the relocation register is added to every address generated by a user process at the time it is sent to memory
 - Base register now called relocation register
 - MS-DOS on Intel 80x86 used 4 relocation registers

Dynamic relocation using a relocation register

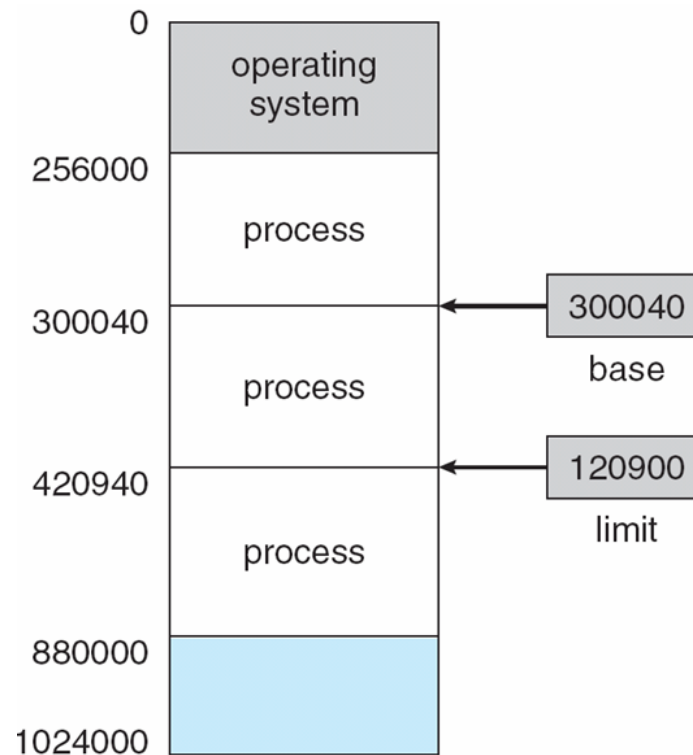


Dynamic allocation using a relocation register

- Routine is not loaded until it is called
 - All routines kept on disk in relocatable load format
 - Better memory-space utilization; unused routine is never loaded
- Useful when large amounts of code are needed to handle infrequently occurring cases
- No special support from the operating system is required
 - Implemented through program design
 - OS can help by providing libraries to implement dynamic loading

HW Protection with Base and Limit Registers

- A pair of base and limit registers define the logical address space
- CPU must check every memory access generated in user mode to be sure it is between base and limit for that user



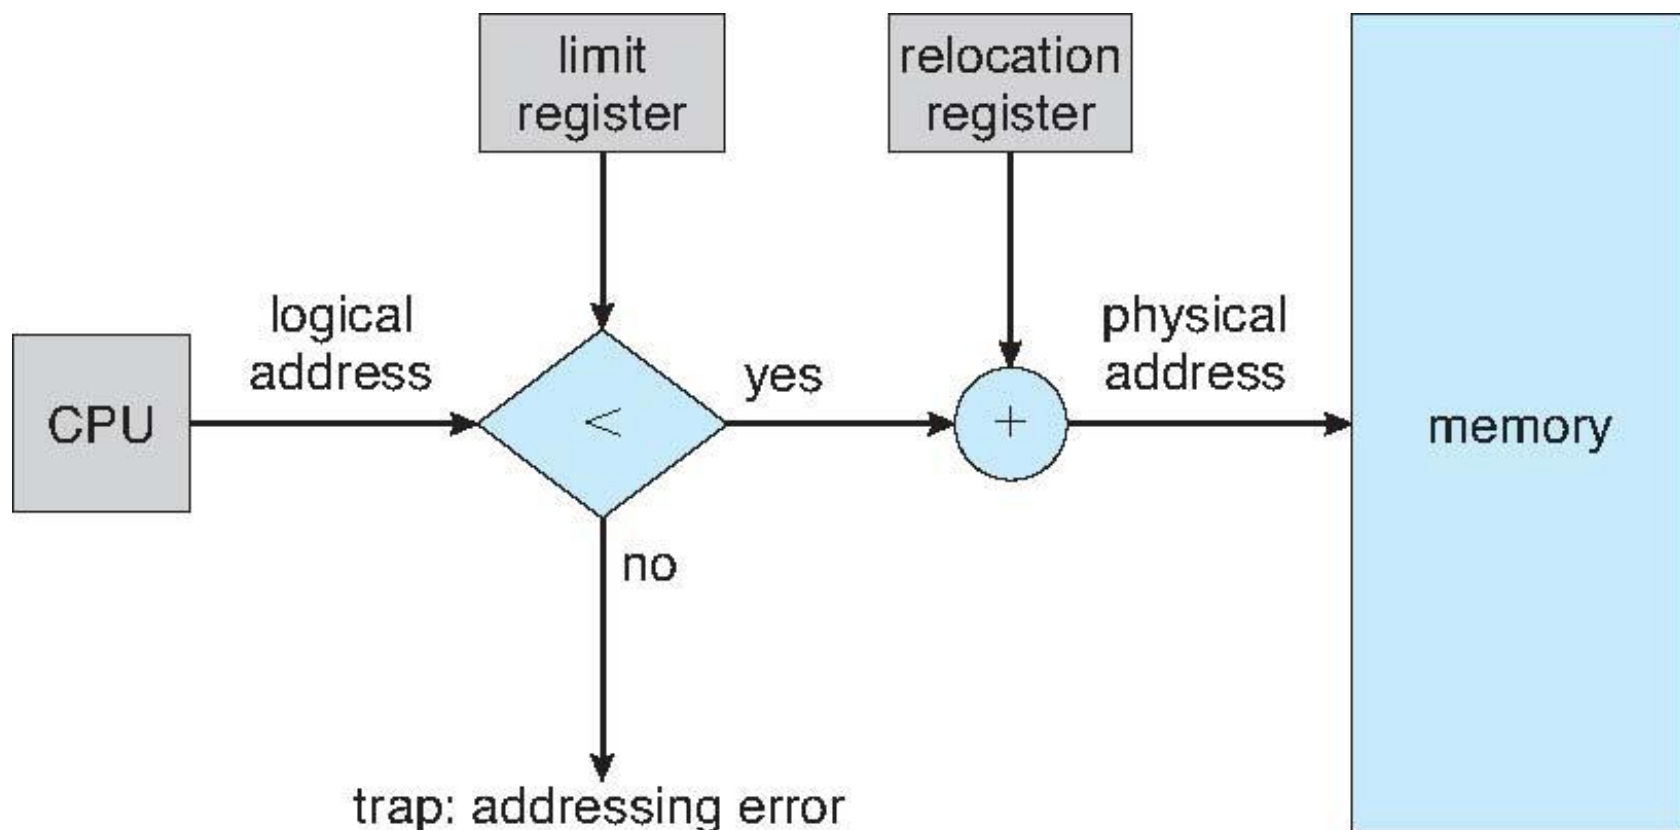
Contiguous Allocation

- Main memory must support both OS and user processes
- Limited resource, must allocate efficiently
- Contiguous allocation is one early method
- Main memory usually into two partitions:
 - Resident operating system, usually held in low memory with interrupt vector
 - User processes then held in high memory
 - Each process contained in single contiguous section of memory

Contiguous Allocation (Cont.)

- Relocation registers used to protect user processes from each other, and from changing operating-system code and data
 - Base register contains value of smallest physical address
 - Limit register contains range of logical addresses – each logical address must be less than the limit register
 - MMU maps logical address dynamically
 - Can then allow actions such as kernel code being transient and kernel changing size

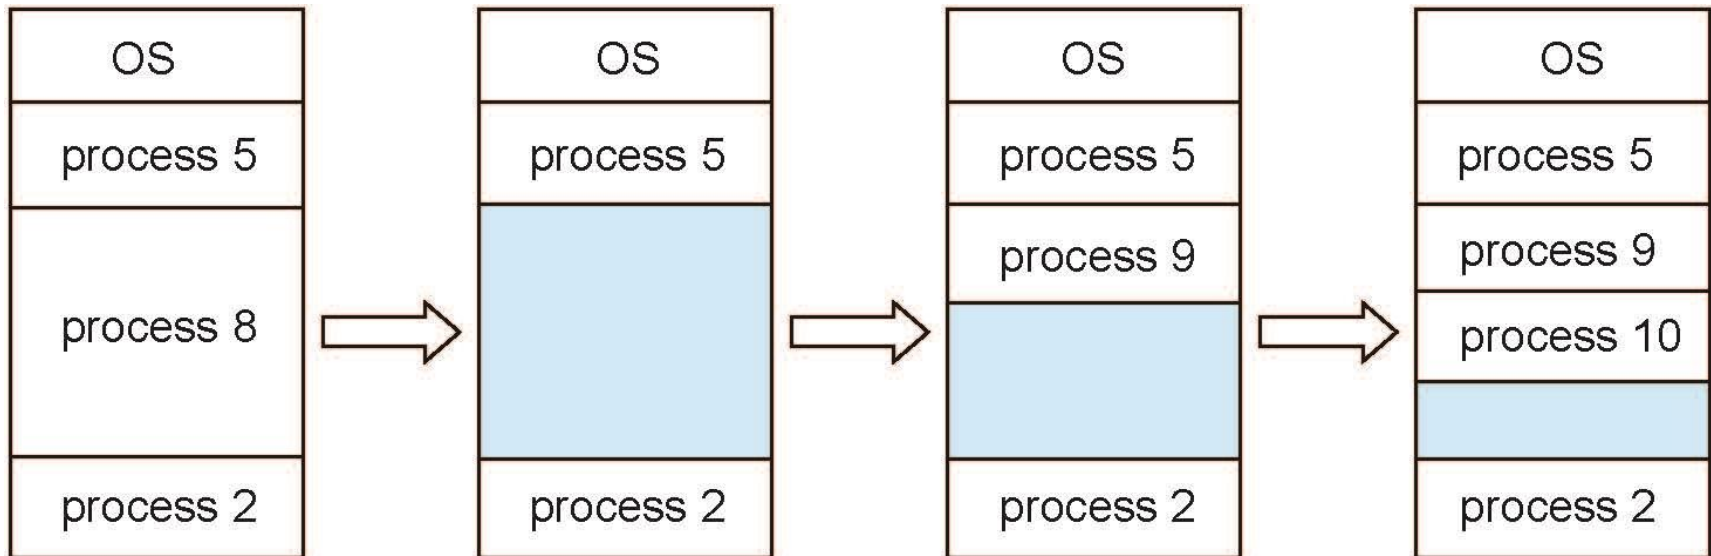
HW Support for Relocation and Limit Registers



Multiple-partition allocation

- Multiple-partition allocation
 - Variable-partition sizes for efficiency (sized to a given process' needs)
 - Hole – block of available memory; holes of various size are scattered throughout memory
 - When a process arrives, it is allocated memory from a hole large enough to accommodate it
 - Process exiting frees its partition, adjacent free partitions combined
 - Operating system maintains information about:
 - a) allocated partitions b) free partitions (hole)

Multiple-partition allocation



Problem of Dynamic Allocation

- How to satisfy a request of size n from a list of free holes?
 - First-fit: Allocate the first hole that is big enough
 - Best-fit: Allocate the smallest hole that is big enough; must search entire list, unless ordered by size
 - Produces the smallest leftover hole
 - Worst-fit: Allocate the largest hole; must also search entire list
 - Produces the largest leftover hole
- First-fit and best-fit better than worst-fit in terms of speed and storage utilization

Issue: Fragmentation

- External Fragmentation – total memory space exists to satisfy a request, but it is not contiguous
- Internal Fragmentation – allocated memory may be slightly larger than requested memory; this size difference is memory internal to a partition, but not being used
- First fit analysis reveals that given N blocks allocated, $0.5 N$ blocks lost to fragmentation
 - $1/3$ may be unusable \rightarrow 50-percent rule


Issue: Fragmentation (Cont.)

- Reduce external fragmentation by compaction
 - Shuffle memory contents to place all free memory together in one large block
 - Compaction is possible only if relocation is dynamic, and is done at execution time
 - I/O problem
 - Latch job in memory while it is involved in I/O
 - Do I/O only into OS buffers
- Now consider that backing store has same fragmentation problems

Issue: Swapping

- A process can be swapped temporarily out of memory to a backing store, and then brought back into memory for continued execution
 - Total physical memory space of processes can exceed physical memory
- Backing store: fast disk large enough to accommodate copies of all memory images for all users; must provide direct access to these memory images
- Roll out, roll in: swapping variant used for priority-based scheduling algorithms; lower-priority process is swapped out so higher-priority process can be loaded and executed
- Major part of swap time is transfer time; total transfer time is directly proportional to the amount of memory swapped
- System maintains a ready queue of ready-to-run processes which have memory images on disk

Issue: Schematic View of Swapping



Issue: Swapping (Cont.)

- Does the swapped out process need to swap back in to same physical addresses?
 - Depends on address binding method
 - Plus consider pending I/O to / from process memory space
- Standard swapping not used in modern operating systems
 - But modified version common
 - Swap only when free memory extremely low
- Modified versions of swapping are found on many systems (i.e., UNIX, Linux, and Windows)
 - Swapping normally disabled
 - Started if more than threshold amount of memory allocated
 - Disabled again once memory demand reduced below threshold

Issue: Context Switch Time including Swapping

- If next processes to be put on CPU is not in memory, need to swap out a process and swap in target process
- Context switch time can then be very high
- 100MB process swapping to hard disk with transfer rate of 50MB/sec
 - Swap out time of 2000 ms
 - Plus swap in of same sized process
 - Total context switch swapping component time of 4000ms (4 seconds)
- Can reduce if reduce size of memory swapped – by knowing how much memory really being used
 - System calls to inform OS of memory use via `request_memory()` and `release_memory()`

Issue: Context Switch Time and Swapping (Cont.)

- Other constraints as well on swapping
 - Pending I/O – can't swap out as I/O would occur to wrong process
 - Or always transfer I/O to kernel space, then to I/O device
 - Known as double buffering, adds overhead

